

Due: 28th February 2020

Total: 50 points

Instructions

- [Institute Plagiarism Policy](#) applies to all HWs.
- It has been made explicit wherever usage of `networkx` or `snap.py` is required. In all other cases, these libraries can only be used for loading and representing the graphs. The implementations need to be done from scratch.
- **Submission Instructions:**
 - All the submissions must be inside a `zip` file named `a1_<your_roll_number>.zip` containing a report named `report.pdf` and a folder named `src` containing all your scripts.
 - All the code must be well documented. Failure to do so will result in a deduction of 2% of the total from the obtained marks in the respective question.
 - All the plots and analyses required should be uploaded in a PDF file named `report.pdf`.

Problem 1: Generating Barabási-Albert Networks

(30 points)

- You can make use of `networkx`, `scipy`, `numpy` etc. for this question
- Generate a network with $N = 10^4$ nodes using the Barabási-Albert model with $m = 4$. Use as initial condition a fully connected network with $m = 4$ nodes.
- Measure the degree distribution at intermediate steps, namely when the network has 10^2 , 10^3 and 10^4 nodes.
- Compare the distributions at these intermediate steps by plotting them together and fitting each to a power-law with degree exponent γ . Do the distributions “converge”? (For fitting the power law, you can take log and perform regression to get the power law parameter, or use any method for fitting a distribution that you find convenient. Libraries are allowed for this question.)
- Plot together the cumulative degree distributions at intermediate steps.
- Measure the average clustering coefficient in function of N .
- Measure the degree dynamics of one of the initial nodes and of the nodes added to the network at time $t = 100$, $t = 1,000$ and $t = 5,000$. For what analysis you have to do, take a look at Image 5.6 at this [link](#).

Problem 2: Implementing PageRank

(20 points)

In this problem, you will learn how to implement the PageRank algorithm. You will be experimenting with a small randomly generated graph (assume graph has no dead-ends) provided in the `graph.txt` file.

It has $n = 100$ nodes (numbered $1, 2, \dots, 100$), and $m = 1024$ edges, 100 of which form a directed cycle (through all the nodes) which ensures that the graph is connected. It is easy to see that the existence of such a cycle ensures that there are no dead ends in the graph. **There may be multiple edges between a pair of nodes, your program should handle these instead of ignoring them.** The first column in `graph.txt` refers to the source node, and the second column refers to the destination node.

Assume the directed graph $G = (V, E)$ has n nodes (numbered $1, 2, \dots, n$) and m edges, all nodes have positive out-degree, and $M = [M_{ji}]_{n \times n}$ is an $n \times n$ matrix as defined in class such that for any $i, j \in [1, n]$:

$$M_{ij} \begin{cases} \frac{1}{\deg(i)}, & \text{if } (i \rightarrow j) \in E \\ 0, & \text{otherwise} \end{cases}$$

Here, $\deg(i)$ is the number of outgoing edges of node i in G . By the definition of PageRank, assuming $1/\beta$ to be the teleport probability, and denoting the PageRank vector by the column vector r , we have the following equation:

$$r = \frac{1 - \beta}{n} \mathbf{1} + \beta M r$$

where, $\mathbf{1}$ denotes $n \times 1$ column vector.

Based on this, the PageRank algorithm runs as follows:

1. Initialise $r^{(0)} = \frac{1}{n} \mathbf{1}$
2. For i from 1 to k , iterate: $r^{(i)} = \frac{1 - \beta}{n} \mathbf{1} + \beta M r^{(i-1)}$

Run the aforementioned iterative process for 40 iterations (assuming $\beta = 0.8$) and obtain the PageRank vector \mathbf{r} . Compute the following:

- List the top 5 node ids with the highest PageRank scores.
- List the bottom 5 node ids with the lowest PageRank scores.